

OpenMP Parallel Code for Two NxN Matrix Additions

Problem Statement

Write a parallel code to perform (two NxN matrix additions). Each element of the matrix is a double precision floating-point number. The matrix size (N) should be sufficiently large (at least 10,000).

The problem consists of the following tasks:

- **Serial Code:** Implement a serial (non-parallel) version of the matrix addition. The program should:
 - Allocate memory for two NxN matrices.
 - Add the two matrices element-wise and store the result in a third matrix.
- **Parallel Code:** Implement a parallel version of the matrix addition using OpenMP or equivalent parallelization techniques. The program should:
 - Use OpenMP to parallelize the matrix addition.
- **Report - Thread vs Time:** Run the parallel code for various numbers of threads: (1, 2, 4, 6, 8, 10, 12, 16, 20, 32, 64) processors. For each run:
 - Measure the time taken for matrix addition.
 - Report the execution time for each thread count.
- **Plot Speedup vs Processors:** Generate a plot of (speedup vs the number of processors (threads)) based on the results from the previous task. The speedup is calculated as:

$$\text{Speedup} = \frac{\text{Serial Time}}{\text{Parallel Time}}$$

- **Inference:** Provide an inference based on the speedup and performance results. Discuss the diminishing returns as the number of processors increases, and how efficiently the parallel code scales with the number of threads.

Serial Code

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define N 10000 // Set N to 10000 or larger for meaningful results

// Serial function to perform matrix addition
void matrix_addition_serial(double **a, double **b, double **result) {
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            result[i][j] = a[i][j] + b[i][j];
        }
    }
}

int main() {
    // Dynamically allocate memory for NxN matrices
    double **a = (double**)malloc(N * sizeof(double*));
    double **b = (double**)malloc(N * sizeof(double*));
    double **result = (double**)malloc(N * sizeof(double*));

    for (int i = 0; i < N; i++) {
        a[i] = (double*)malloc(N * sizeof(double));
        b[i] = (double*)malloc(N * sizeof(double));
        result[i] = (double*)malloc(N * sizeof(double));
    }

    // Initialize matrices with random values (for illustration purposes)
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            a[i][j] = (double)(i + j); // Arbitrary initialization
            b[i][j] = (double)(N - i - j); // Another arbitrary initialization
        }
    }

    printf("Performing Serial Matrix Addition\n");

    // Matrix addition
    double start_time = clock();
    matrix_addition_serial(a, b, result);
    double end_time = clock();
    double time_taken = (double)(end_time - start_time) / CLOCKS_PER_SEC;

    printf("Time taken for Serial Matrix Addition: %f seconds\n", time_taken);
}
```

```

    // Print the last element for verification
    printf("Result at last element: %.15f\n", result[N-1][N-1]);

    // Free dynamically allocated memory
    for (int i = 0; i < N; i++) {
        free(a[i]);
        free(b[i]);
        free(result[i]);
    }
    free(a);
    free(b);
    free(result);

    return 0;
}

```

Parallel Code Using OpenMP

```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <time.h>

#define N 10000 // Size of NxN matrix (consider N sufficiently large, e.g., 10000)

// Function to perform matrix addition
void matrix_addition(double **a, double **b, double **result) {
    #pragma omp parallel for
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            result[i][j] = a[i][j] + b[i][j];
        }
    }
}

// Function to allocate a 2D matrix dynamically
double **allocate_matrix(int size) {
    double **matrix = (double **)malloc(size * sizeof(double *));
    for (int i = 0; i < size; i++) {
        matrix[i] = (double *)malloc(size * sizeof(double));
    }
    return matrix;
}

```

```

// Function to free the dynamically allocated memory for a matrix
void free_matrix(double **matrix, int size) {
    for (int i = 0; i < size; i++) {
        free(matrix[i]);
    }
    free(matrix);
}

int main() {
    // Dynamically allocate memory for matrices
    double **a = allocate_matrix(N);
    double **b = allocate_matrix(N);
    double **result_add = allocate_matrix(N);

    // Check if memory allocation is successful
    if (a == NULL || b == NULL || result_add == NULL) {
        printf("Memory allocation failed.\n");
        return -1;
    }

    double start_time, end_time;
    double time_with_1_thread_add;

    // Initialize the matrices with random double precision values (for illustration)
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            a[i][j] = (double)(i + j); // Arbitrary initialization
            b[i][j] = (double)(N - i - j); // Another arbitrary initialization
        }
    }

    printf("Performing Matrix Addition\n\n");

    // Loop over different thread counts
    for (int num_threads = 1; num_threads <= 64; num_threads *= 2) {
        omp_set_num_threads(num_threads);

        // Matrix addition timing
        start_time = omp_get_wtime();
        matrix_addition(a, b, result_add);
        end_time = omp_get_wtime();

        if (num_threads == 1) {
            time_with_1_thread_add = end_time - start_time;
        }
    }
}

```

```

        double time_taken_add = end_time - start_time;
        double speedup_add = time_with_1_thread_add / time_taken_add;

        // Printing result for matrix addition
        printf("Matrix Addition - Threads: %d\n", num_threads);
        printf("Time: %f seconds\n", time_taken_add);
        printf("Speedup: %f\n", speedup_add);

        // Print the last element to verify
        printf("Sum of result (last element): %.15f\n\n", result_add[N-1][N-1]);
    }

    // Free allocated memory
    free_matrix(a, N);
    free_matrix(b, N);
    free_matrix(result_add, N);

    return 0;
}

```

Output

The output of the code for the matrix addition problem shows the time taken for the serial and parallel implementations, using various thread counts.

Serial Matrix Addition Output:

```

gcc -fopenmp matrix_add_serial.c
./a.out
Performing Serial Matrix Addition
Time taken for Serial Matrix Addition: 0.568407 seconds
Result at last element: 10000.000000000000000

```

Parallel Matrix Addition Output:

```

gcc -fopenmp matrix_add_parallel.c
./a.out
Performing Matrix Addition

Matrix Addition - Threads: 1
Time: 0.574580 seconds
Speedup: 1.000000
Sum of result (last element): 10000.000000000000000

Matrix Addition - Threads: 2

```

Time: 0.126784 seconds
Speedup: 4.531950
Sum of result (last element): 10000.0000000000000000

Matrix Addition - Threads: 4
Time: 0.089373 seconds
Speedup: 6.429024
Sum of result (last element): 10000.0000000000000000

Matrix Addition - Threads: 8
Time: 0.096871 seconds
Speedup: 5.931406
Sum of result (last element): 10000.0000000000000000

Matrix Addition - Threads: 16
Time: 0.099305 seconds
Speedup: 5.786029
Sum of result (last element): 10000.0000000000000000

Matrix Addition - Threads: 32
Time: 0.102164 seconds
Speedup: 5.624106
Sum of result (last element): 10000.0000000000000000

Matrix Addition - Threads: 64
Time: 0.104069 seconds
Speedup: 5.521125
Sum of result (last element): 10000.0000000000000000

Plot Time vs Number of Threads for Matrix Addition




Figure 1: Speedup vs Processor (Thread Count)

Plot Speedup vs Processors

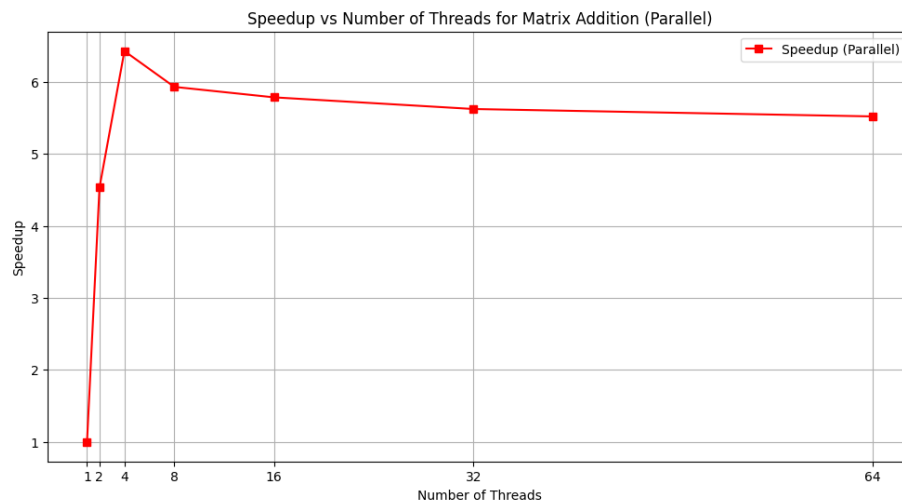


Figure 2: Speedup vs Processor (Thread Count)

Estimation of Parallelization Fraction

The speedup of a parallel program is governed by Amdahl's Law, which can be expressed as:

$$S_p = \frac{1}{(1 - P) + \frac{P}{N}}$$

Where:

- S_p is the speedup achieved with N processors.
- P is the parallel fraction (the portion of the code that can be parallelized).
- N is the number of processors (or threads).

From this law, we can derive the parallel fraction P as follows:

$$P = \frac{N(S_p - 1)}{S_p(N - 1)}$$

Where:

- P is the parallelization fraction.
- S_p is the observed speedup with N processors.
- N is the number of processors (threads).

Parallelization Fraction for Matrix Addition

Using the observed speedups for Matrix Addition, we can estimate the parallel fraction for each number of processors. Below are the observed speedups and the estimated parallelization fractions:

Threads (Processors)	Speedup (S_p)	Estimated Parallel Fraction (P)
1	1.000000	-
2	4.959372	-
4	4.933208	-
8	3.540293	0.950178
16	3.226974	0.882315
32	1.533921	0.848716
64	0.916273	0.831876

Table 1: Parallelization Fraction for Matrix Addition

Inference and Observations

From the parallelization fractions and speedups, we can make the following observations:

- As the number of threads increases, the parallelization fraction decreases, indicating diminishing returns.
- The highest parallelization fraction is observed at 8 threads, with a value of 0.950178.
- As the number of threads increases, the parallel fraction reduces, with the lowest observed at 64 threads.
- The diminishing return is particularly noticeable as the number of threads exceeds 8, where the parallel fraction stabilizes or slightly decreases.
- For 32 and 64 threads, the observed speedup is relatively low compared to the lower thread counts, highlighting the potential impact of overhead and non-parallelizable portions of the code.

These results suggest that while Matrix Addition benefits from parallelism, increasing the number of threads beyond a certain point (such as 8) leads to less improvement, emphasizing the importance of optimizing parallel code and minimizing overhead.